

04-06-2026

**How to build a real-life theory
of change from your
interviews using AI:
a Causal Map 4 training
session**

The team


Steve Powell, Co-founder and Director,
Causal Map Ltd



Gabriele Caldas Cabral, Outreach
Coordinator, Causal Map Ltd



What we will do today: gather interviews with AI, turn them into a theory of change, and learn to question it. About 90 minutes, with plenty of room for your questions.




You have a stack of interviews

Somewhere in there is your answer

Interviews. Reports. Open survey answers. Pages of people telling you what changed and why.

Somewhere in there is the answer to your question. How do you get it out, in a way you can stand behind?




You have a stack of interviews

Two tempting shortcuts, both bad

Hand it to the black box. “ChatGPT, what does this say?” Fast and fluent, but you cannot see what it leaned on, so you cannot defend the answer.

Read it all yourself. Thorough, but it does not scale. Three hundred transcripts, and you still have to show your synthesis really reflects what people said.

There is a better way.




The plan for today

One route from your interviews to a theory of change you can defend:


- start from interviews, reports, open answers, any narrative data
- turn what people said into a map of causes and effects, with AI
- question that map until you have an answer

AI does the patient clerical work. You make the judgements.



You have a stack of interviews

Why causal mapping




What we do

You read the text and mark each causal claim as a link from one factor to another.

"The training gave me confidence, and that is why I started the business."

This becomes **training** → **confidence** → **started a business**, keeping the exact quote and the source on every link.

Do that across all your interviews and the links join up into a map: your theory of change, built from what people actually said.



Non-causal Qualitative Data Analysis (QDA) version 1:
code causality as a theme/tag/concept

Text and quote	Theme / tag
"the floods destroyed our crops"	Floods destroy crops (causal)

Non-causal QDA version 2: code themes/tags/concepts,
with optional post-hoc connections between themes

Text and quote	Theme / tag
"the floods destroyed our crops"	Floods
"the floods destroyed our crops"	Crops destroyed


Optional post-hoc
connection

Causal QDA: code causal claims natively between *pairs* of codes



Text and quote	"Cause"	"Effect"
"the floods destroyed our crops"	Floods happened	Crops destroyed


Quote	Cause	Effect	Source ID
The floods destroyed our crops	Floods happened	Crops destroyed	A3
the chickens died too	Floods happened	Chickens died	A3
Our harvest was ruined due to the high water	Floods happened	Crops destroyed	A4



We code claims, not facts

A link means: this person says X influenced Y. Not that X really did.


Twenty people saying so is not proof. It is evidence you can now weigh. Turning that evidence into a conclusion is your job, and we come back to it later.



Keep it simple

We do not code how strong the link was, or whether it was good or bad, or what the hidden meaning might be.

People say “X made Y happen”. They rarely say how strongly. So we do not invent it. Simple coding stays close to what people said, and that is exactly why AI can do it.




Why use AI for causal mapping

AI as a clerk, not an oracle

The clerk's job: find every causal claim, attach a quote. Tireless, exhaustive, cheap.

Your job: decide the question, check the work, judge what it all means.



The clerical task is simple enough to hand over. The judgement is not, so you keep it.



Built for scale


A team reading by hand tires, gets bored, and pays uneven attention. The AI does not.


It codes every claim across hundreds of documents, and each one traces back to the sentence it came from. So you capture everything cheaply, then spend that volume down into a conclusion you can defend.




Case studies

- **VIB, Brussels:** 100 HR interviews on career trajectories
- **Sudan:** 60 development professionals for a programme evaluation
- **Gender gap in STEM, Chile:** 32 student interviews
- **UK postdoc feedback:** mid-career drivers and obstacles to learning
- **Climate adaptation, Austria:** 38 harvested outcome statements








INTRAC needed the big picture from a programme across 13 countries. The AI coded 5,430 causal claims into an overall map and one per country, and drafted the written analysis as a vignette.


The workflow, three steps



Step 1: Start from the question

Before you touch the data, write down what you want to be able to say at the end, and to whom.

Every later choice follows from that one sentence: the data you gather, the labels you use, the maps you make.



Be realistic about what it can answer


Good at

- which factors matter most
- what drives or follows from a factor
- how groups differ
- whether the evidence fits your theory of change

Pick questions the method can serve.

Not for

- effect sizes
- proving X causes Y on its own




Step 2: Code the claims

You write a short instruction for the AI, like a chatbot prompt, and it codes the links for you.

The golden rule: test on a small, varied sample, see exactly where the output is wrong or thin, fix the instruction, then scale up.

And one rule you never break: **every link needs a quote**. Without it you cannot show your working.




Step 3: Check the links

However careful the coding, some links will be wrong. Check them before you analyse.

- tag the doubtful or surprising ones, to filter later
- note how sure each source sounds
- note how reliable each source is

A quick read of a varied sample catches most problems.




**Your map is a theory of change
you can question**

A filter is a question

Your links are not a static report. They are a living map you can question, over and over.

You ask a question by filtering. Women only. Everything downstream of training. Stack the filters and you answer a bigger question. The same data gives very different maps, each one just the result of a different question.




Your map is a theory of change you can question

Two everyday examples

“What did the cash transfer lead to?” Filter to that factor, look downstream. Out comes a map of every reported consequence, with counts.

“Do women and men tell different stories?” Split the same map by group. The links each group stresses light up differently.




The transitivity trap

A pig farmer says the cash grant gave them **more cash**. A wheat farmer says **more cash** let them buy **more seed**.

So cash grants lead to more seed? **No**. Two people, two stories, stitched into one that nobody told.


The safe move: keep only the sources whose own account runs all the way through.




Step back and judge

Behind one tidy map there may still be hundreds of quotes. Does the claim hold up? Do the links really belong to the same context?


The AI can draft a source-by-source commentary on each pathway, doing only what a patient reader could. Treat it as a first draft and edit it.








A real evaluation: the Love Alliance



Key populations are healthy and empowered through inclusive people-centred SRHR

We build on past achievements



The official theory of change: a tidy ladder from strategies up to goals. Clear, communicable, fundable. But a ladder cannot draw a loop.

A real evaluation: the Love Alliance


Too much evidence, the nicer problem

An end-term evaluation across ten African countries, with Southern Hemisphere. We coded the narrative data for them.


- 176 documents: 79 country reports, 97 interview transcripts
- over 22,000 causal claims, coded by AI
- 13,756 kept for the maps

No team could hand-code that and stay fresh. Every link still carries the quote it came from.





A map built from the interviews. The biggest factor, Love Alliance support for advocacy and capacity building, was cited 777 times. The story people told was “we built each other up”, which no ladder predicted.



The movement-building view. Networking is the hub, with reinforcing circles: capacity feeds peer support feeds capacity; networking enables advocacy which feeds more networking.

Put a doubt to the data


Partway through, a fair worry surfaced: was the programme's advocacy provoking a backlash against key populations?

Rather than trade impressions, we put the question to the data and traced it through the 43 sources that raised it. Every account pointed the other way: the support helped communities resist backlash, none named it as the cause.

"also helping to counter the backlash against women's rights and LGBT rights"




So what



The whole thing in one line

Code “X influenced Y, with a quote”. Capture everything, capture cheaply, let AI do it. Then judge, in the open.


This is not statistical proof. It is a disciplined way to assemble evidence, weigh it transparently, and reach a conclusion you can stand behind.



Need the data first?

No interviews yet? Qualia can gather them. It runs a chat interview that asks people what changed and why, follows up in their own words, in any language.


The answers come back as narrative text, ready to code into a map.



Come and try it

This is how we work at Causal Map Ltd, every day.

If you want to go from a stack of interviews to a theory of change you can defend, come and try it with us at **app.causalmap.app**.





global
evaluation
initiative

Obrigada!

www.glocalevalweek.org

So what